




# Data-Driven Night

---

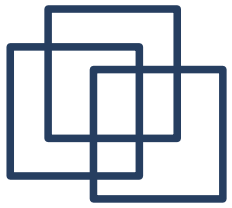
Thessaloniki Java Meetup Group





---


# GENERAL




# What is C\* (I)

---

- Fast Distributed DB
- High Availability
- Near-Linear Horizontal Scalability

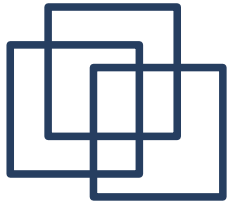




## What is C\* (II)

---


- Predictable Performance
- Fault Tolerance (P2P)
- Cannot replace RDBMS ad hoc
  - Data Model is different!
  - Transactions
    - Kind of native in RDBMS
    - C\* lightweight transaction mechanism exist but should be generally avoided
    - Transactions can be implemented in the Application




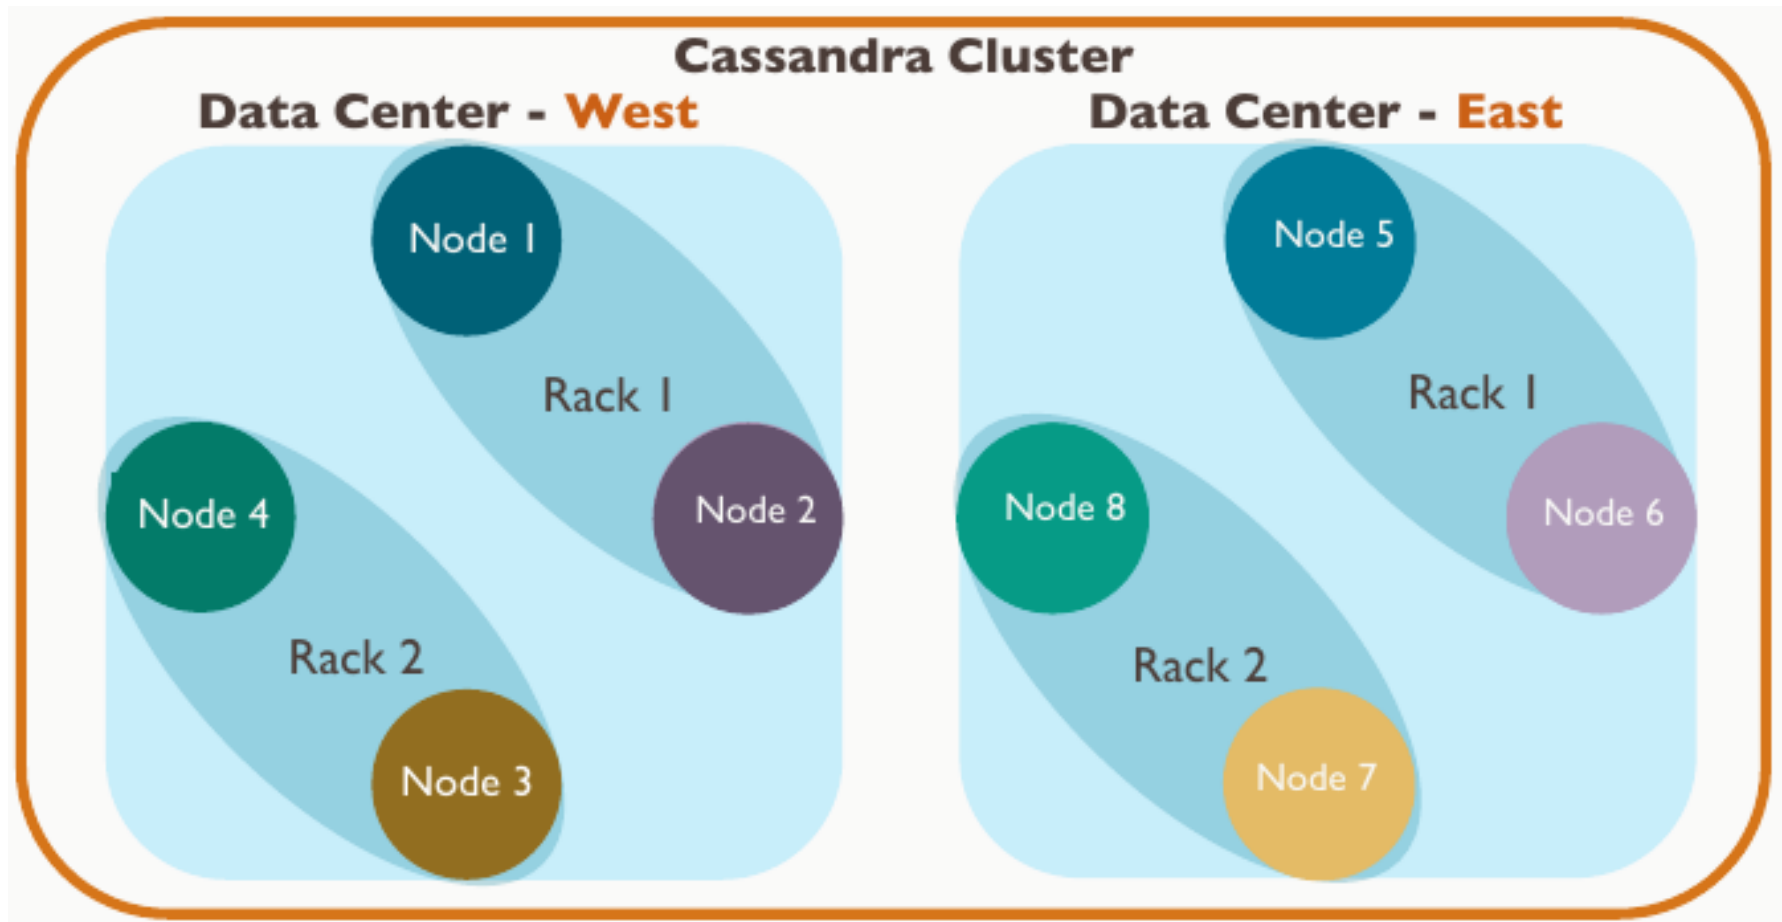
# Origins


---

- Google Big Table
  - Storage Model
- Amazon Dynamo
  - Distribution backbone
- Facebook integrated these two (2008)
  - Later released as Cassandra
  - Nowadays an Apache project




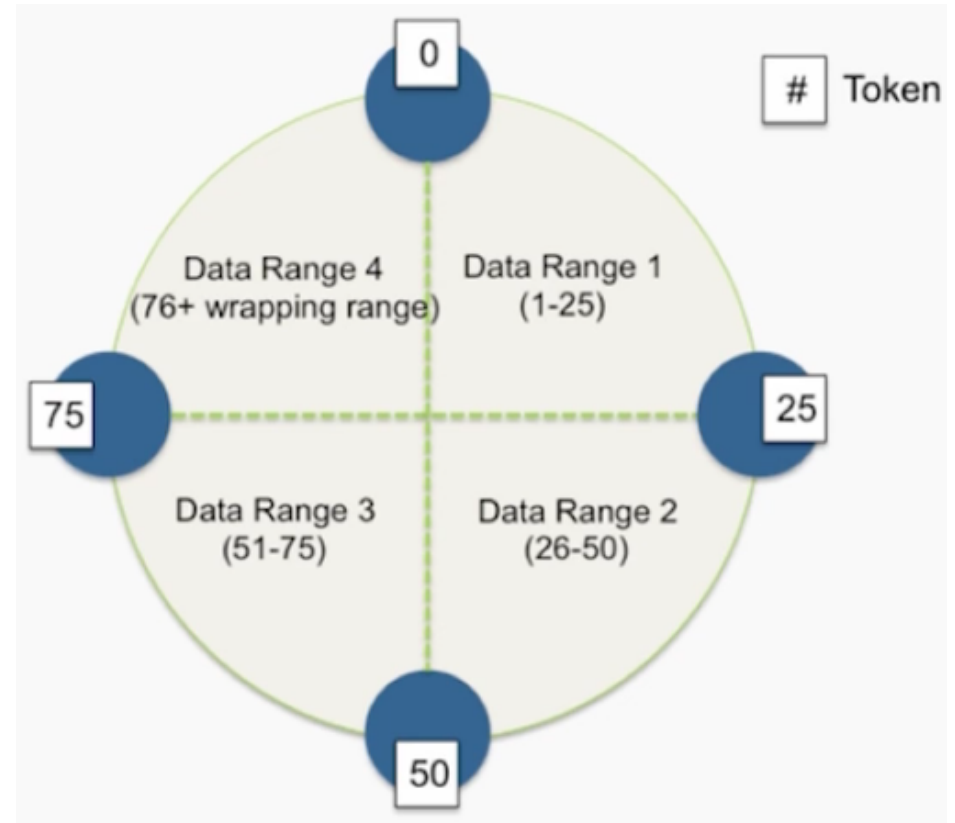
# Structure (I)

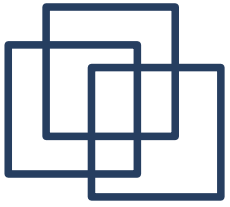




## Structure (II)


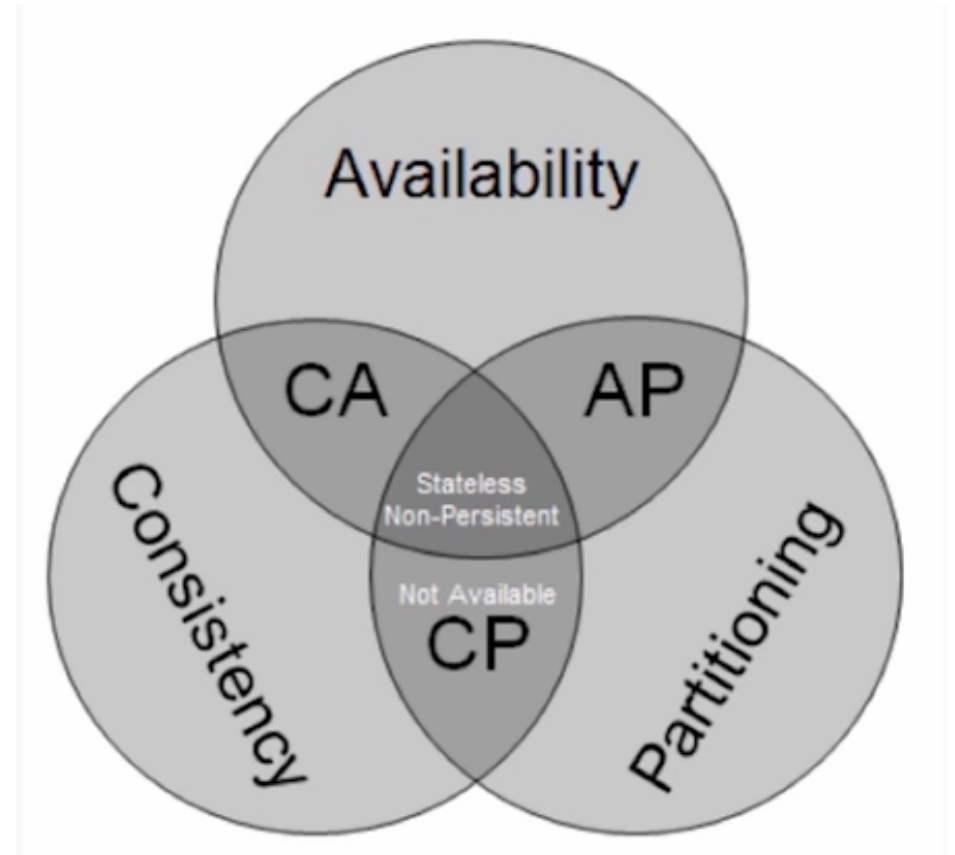
- Hash Ring
- P2P
- Data partitioning
- Replication across peers






# CAP theorem (I)


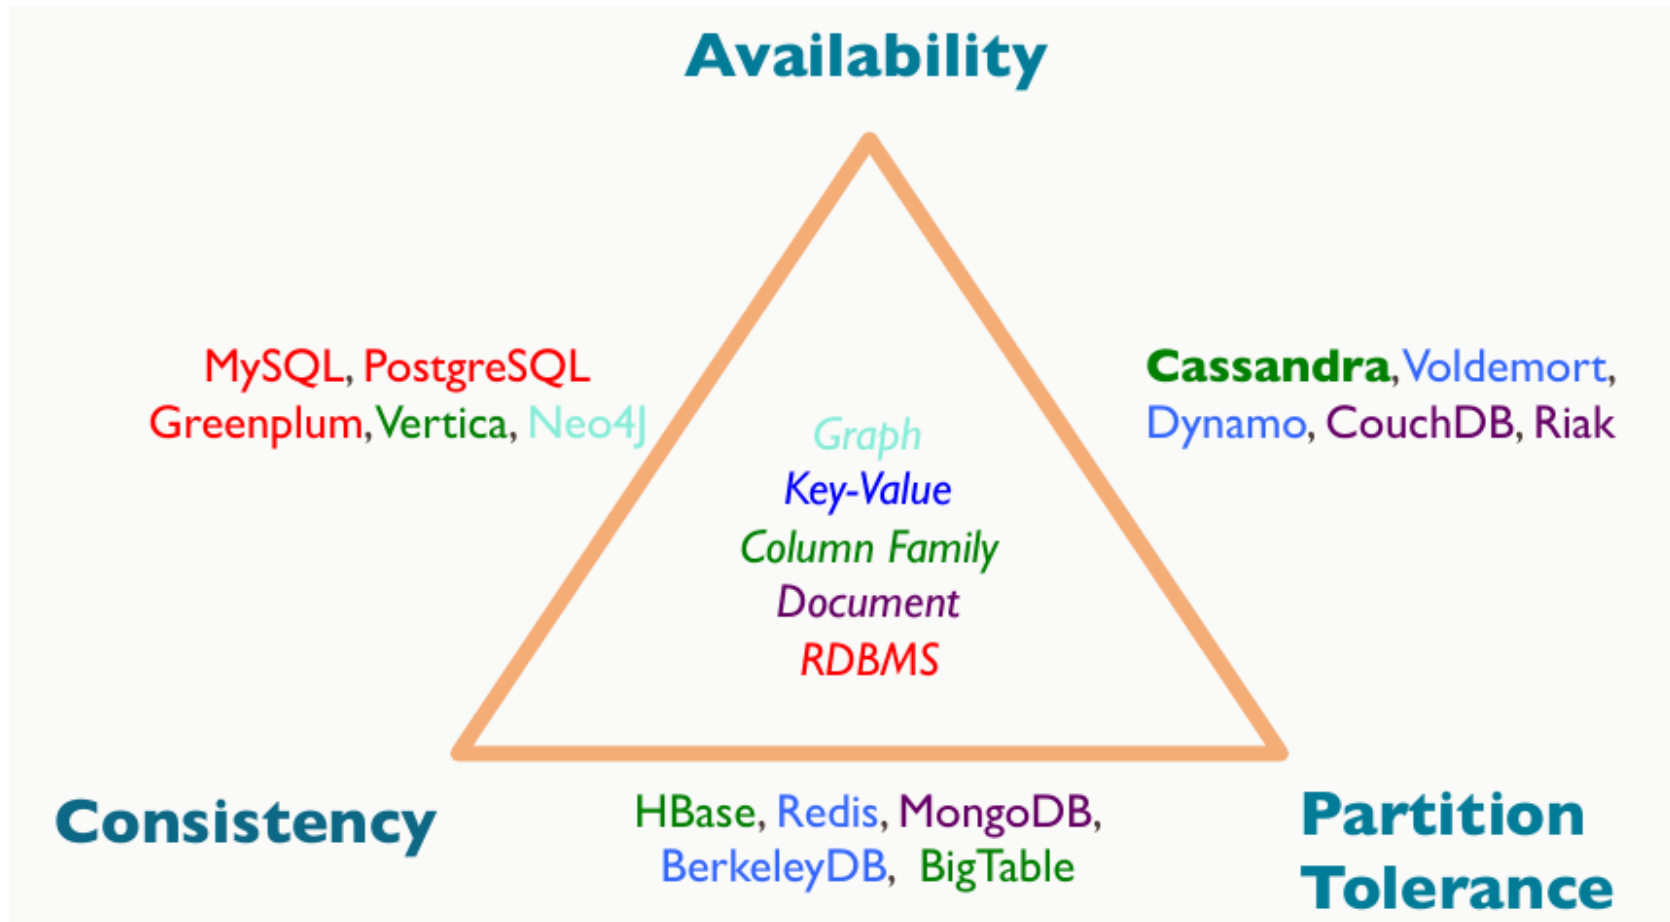
- Consistency  
Availability  
Partitioning trade-off
- Partitioning =  
Partition Tolerance =  
same network or not
- C\* design choice:
  - A, P over C

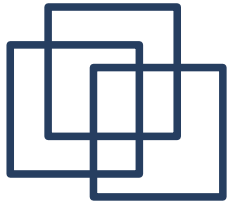






# CAP theorem (II)






# Common Use Cases

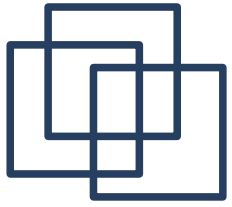
---

- From:  
<http://www.planetcassandra.org/apache-cassandra-use-cases/>
- Product Catalog and Playlist
- Recommendation and Personalization
- Fraud Detection
- Messaging
- IOT and Sensor Data
- Marketing and Advertising
- Social Media and Networking



---

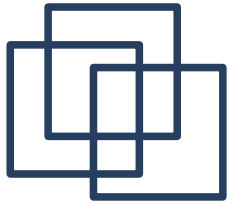
# INTERNALS



# Replication Factor

---

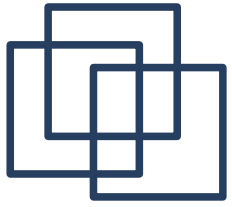
- Definition “how many copies of our data, do exist in a cluster” (RF)
- Data is always replicated
- RF is defined and configured for a KeySpace per Data Center
- A KeySpace is a “collection of Tables”



# Multiple Data Centers

---


- DCs can be physical or logical
- Asynchronous replication to other Dcs
- ```
CREATE KEYSPACE hospital  
WITH REPLICATION = {  
    'class' : 'SimpleStrategy',  
    'replication_factor' : 3  
};
```



# Consistency Level

---

- Definition “How many replicas respond Properly to a query” in order to consider the query successful
  - A query can be a Read or a Write
- Examples: ALL, QUORUM, ONE
- Consistency Level (CL) affects performance and availability (fault-tolerance)
- CL is configured per query
  - This enables using C\* even in **CAP** mode




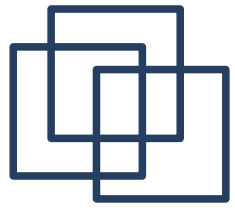
# Consistency Level Details

---

- Several are available
- Defined per request, by default ONE

| Name                     | Description                                                               | Usage                                                |
|--------------------------|---------------------------------------------------------------------------|------------------------------------------------------|
| <b>ANY</b> (writes only) | Write to any node, and store <i>hinted handoff</i> if all nodes are down. | Highest availability and lowest consistency (writes) |
| <b>ALL</b>               | Check all nodes. Fail if any is down.                                     | Highest consistency and lowest availability          |
| <b>ONE</b> (TWO,THREE)   | Check closest node to coordinator.                                        | Highest availability and lowest consistency (reads)  |
| <b>QUORUM</b>            | Check quorum of available nodes.                                          | Balanced consistency and availability                |



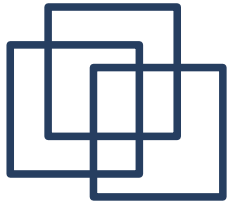


# Consistency Level Trade-Off

---

- Consistency Level ALL
  - **Consistent Read**,  
Highest latency, Lowest availability
- Consistency Level ONE
  - Maybe inconsistent Read,  
Lowest latency, **Highest availability**
- Consistency Level QUORUM
  - Consistent Read (if both Read/Write are QUORUM),  
Medium latency, Medium availability

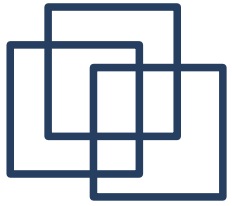




# Immediate Consistency (I)

---


- Immediate Consistency
  - Reads always return the most recent data
- We achieve this by configuring
  - CL per Read, Write
  - RF per KeySpace
- It must hold:  $CL_{Read} + CL_{Write} > RF$
- Practically, does it worth it?
  - CL ONE is enough in most cases



## Immediate Consistency (II)

---


- Configuration examples for a Cluster with 4 Nodes
- Frequent Read operations:
  - $RF = 3$
  - $CL_{Read} = QUORUM, CL_{Write} = QUORUM$
- Frequent Write operations:
  - $RF = 3$
  - $CL_{Read} = ALL, CL_{Write} = ONE$



# Cluster internal communication



---


- Nodes continuously communicate and exchange information
- Two central mechanisms
  - Gossip
  - Snitch



# Gossip

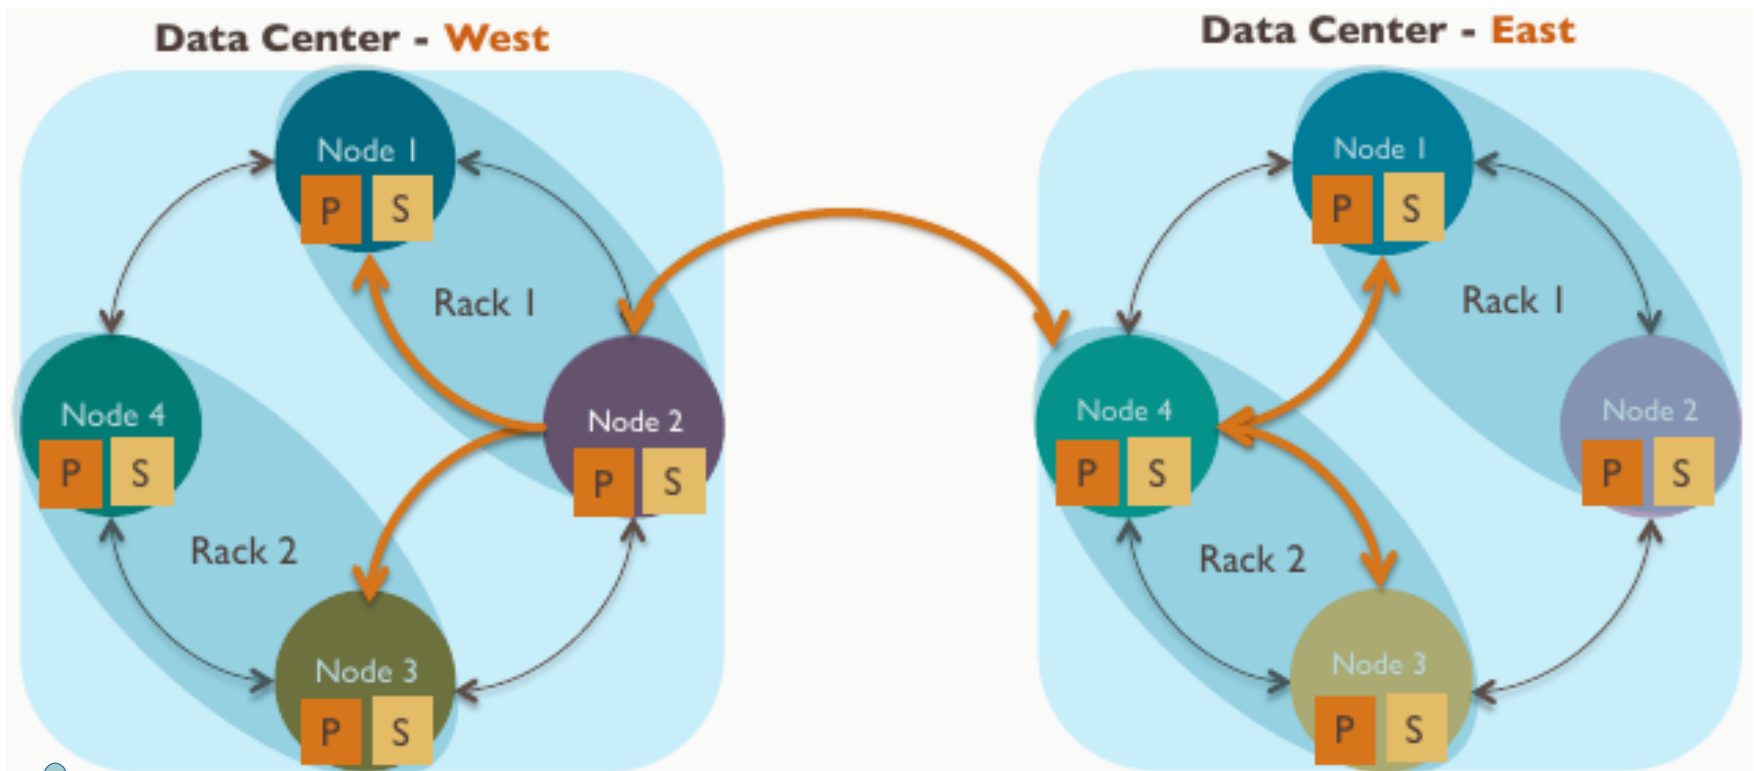
- Every one second, each Node contacts 1 to 3 others, sending and requesting timestamped updates about known Nodes
  - states
  - locations






# Snitch


- This is how Nodes know about the rack and data center topology





---


# CONFIGURATION FILES & TOOLS



# Installation

---

- Requirements for CPU, RAM, HDD
- Operating System
- NTP – C\* requires synchronized clocks
- Disable memory swaps
- Java: Oracle JDK
- Network configuration
- C\* installation
- C\* configuration

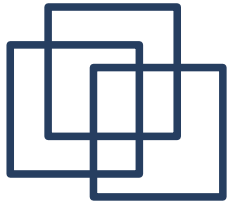


# Distributions

---

- Apache Cassandra
- DataStax Community Edition (DSC)
  - Additional tools for managing a Cluster
- DataStax Enterprise Edition (DSE)
  - More features than DSC, better for Analytics
  - Special program for start-ups
- [www.planetcassandra.org/cassandra/](http://www.planetcassandra.org/cassandra/)






# Configuration Files

---


- Located under **\$CASSANDRA\_HOME/conf/**
  - Example: dsc-cassandra-2.1.10/conf/
- Most important files:
  - cassandra.yaml
  - cassandra-env.sh
  - logback.xml
  - cassandra-rackdc.properties
  - cassandra-topology.properties



# C\* Tools


---

- Located under
  - **`$CASSANDRA_HOME/bin/`**
  - **`$CASSANDRA_HOME/tools/`**
- Tools
  - nodetool
  - cqlsh
  - cassandra-stress
  - sstable2json, json2sstable
  - Cassandra Cluster Management – CCM (DataStax)
  - DevCenter (DataStax)



---


# C\* DATA MODEL & CQL



## C\* Data Model (I)

---

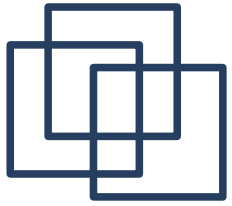
- Data is stored and organized in a Column Family
- A Column Family is comprised of Rows
- A Row is the smallest unit that stores related data



## C\* Data Model (II)

---


- A Partition (old name: RowKey) uniquely identifies a Row in a Column Family
  - It stores data in Cells
  - Cell parts
    - column name
    - column value
    - data creation timestamp
  - Maximum cell size (column value)
    - 2 GB in theory
    - 100MB in praxis



## C\* Data Model (III)

---


- A Table is a 2D view of a column family
  - A table has Partitions
  - A Partition may be a single row or multiple rows
- A Partition Key uniquely identifies a Partition
  - Can be composite
  - It is hashed by the partitioner system to determine which Node will store it



## C\* Data Model (IV)

---

- A Primary Key uniquely identifies a row
  - Can be composite
  - It is comprised of two parts
    - the Partition Key
    - optionally, further columns
- Data Definition Language (DDL) describes Tables, Partition Keys, Primary Keys



# Data in Clustering Columns (I)


---

- For table Videos below:

```
CREATE TABLE Videos (  
    id INT, name TEXT, year int, runtime int,  
    PRIMARY KEY ((year), name)  
);
```

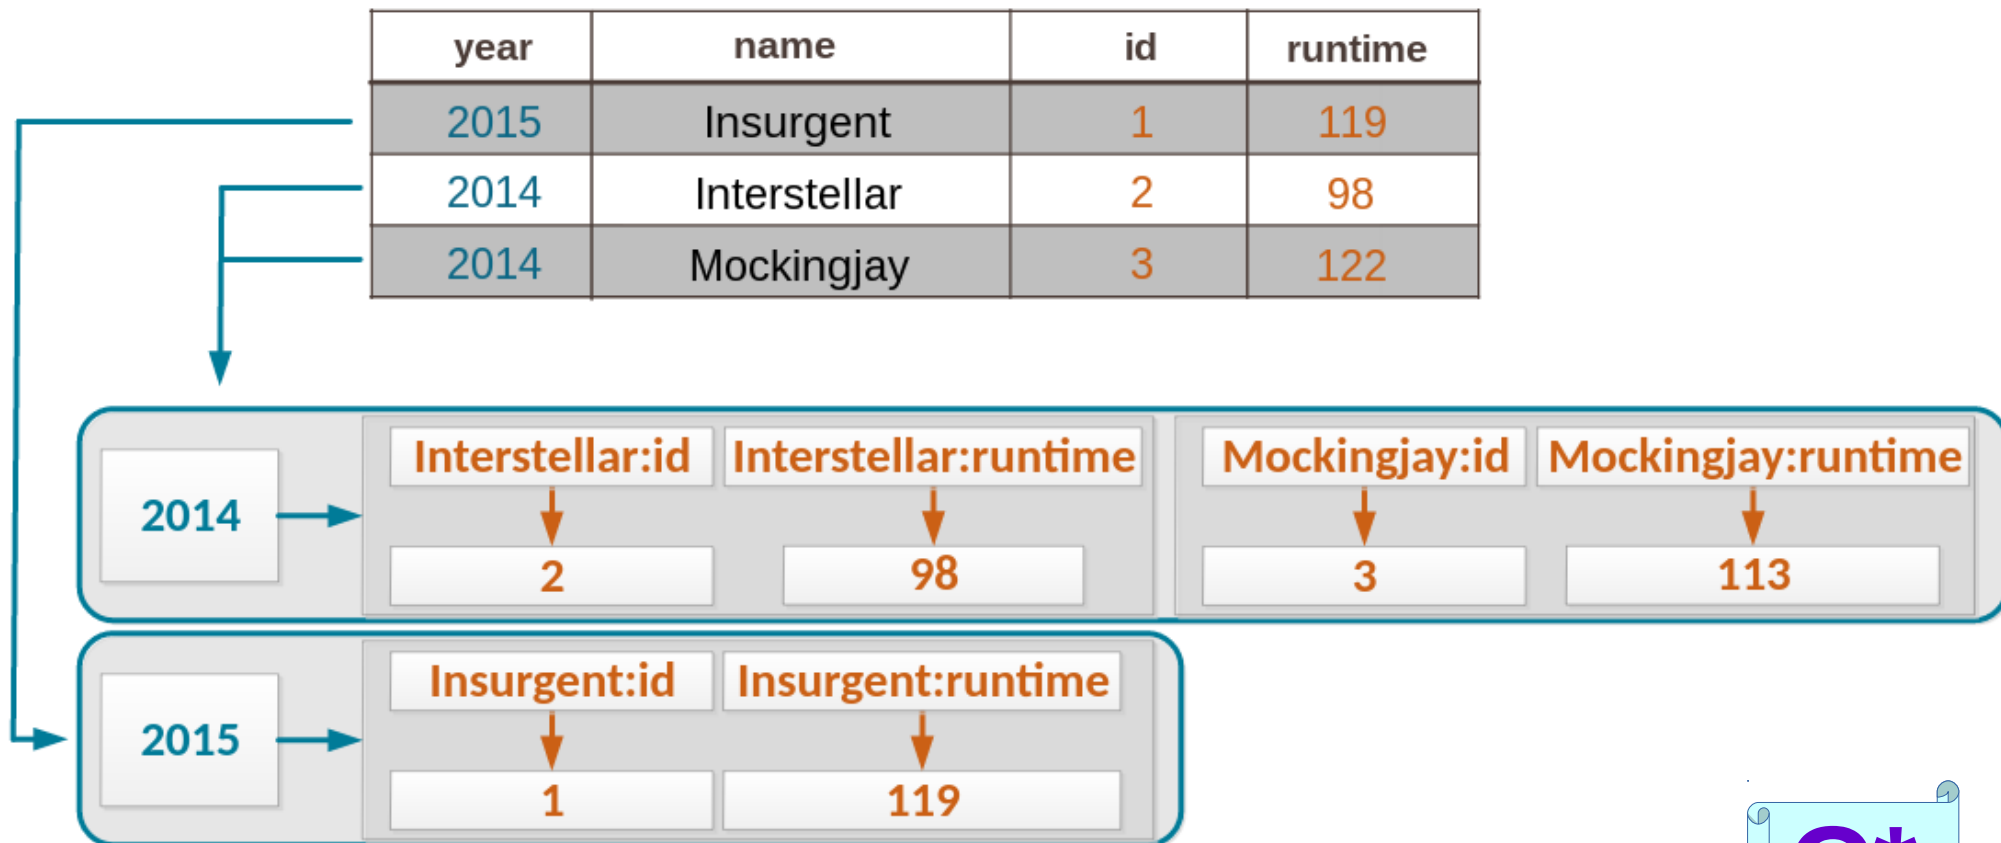
| year | name         | id | runtime |
|------|--------------|----|---------|
| 2015 | Insurgent    | 1  | 119     |
| 2014 | Interstellar | 2  | 98      |
| 2014 | Mockingjay   | 3  | 122     |

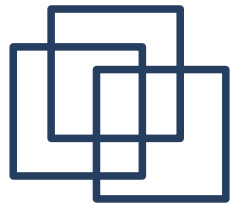




## Data in Clustering Columns (II)

- Clustering columns divide Rows among partitions





# Cassandra Query Language


---


- Language for communicating with the C\* DB
- Abbreviated as CQL
- Similar to SQL
- It can create, modify, delete tables and data



# CQL Basic Data Types

| CQL Type  | Constants         | Description                                                                  |
|-----------|-------------------|------------------------------------------------------------------------------|
| ASCII     | strings           | US-ASCII character string                                                    |
| BIGINT    | integers          | 64-bit signed long                                                           |
| BLOB      | blobs             | Arbitrary bytes (no validation), expressed as hexadecimal                    |
| BOOLEAN   | booleans          | true or false                                                                |
| COUNTER   | integers          | Distributed counter value (64-bit long)                                      |
| DECIMAL   | integers, floats  | Variable-precision decimal                                                   |
| DOUBLE    | integers          | 64-bit IEEE-754 floating point                                               |
| FLOAT     | integers, floats  | 32-bit IEEE-754 floating point                                               |
| INET      | strings           | IP address string in IPv4 or IPv6 format*                                    |
| INT       | integers          | 32-bit signed integer                                                        |
| LIST      | n/a               | A collection of one or more ordered elements                                 |
| MAP       | n/a               | A JSON-style array of literals: { literal : literal, literal : literal ... } |
| SET       | n/a               | A collection of one or more elements                                         |
| TEXT      | strings           | UTF-8 encoded string                                                         |
| TIMESTAMP | integers, strings | Date plus time, encoded as 8 bytes since epoch                               |
| TUPLE     | n/a               | Up to 32k fields                                                             |
| UUID      | uuids             | A UUID in standard UUID format                                               |
| TIMEUUID  | uuids             | Type 1 UUID only (CQL 3)                                                     |
| VARCHAR   | strings           | UTF-8 encoded string                                                         |
| VARINT    | integers          | Arbitrary-precision integer                                                  |






# CQL: Create Table

---


```
CREATE TABLE cars_by_cost (  
    brand TEXT,          # part of Partition Key  
    model TEXT,          # part of Partition Key  
    cost DECIMAL,        # Clustering Key  
    merchant TEXT,  
    PRIMARY KEY ((brand, model), cost)  
) WITH CLUSTERING ORDER BY (cost ASC);
```



## CQL: Modify Table

---

- `ALTER TABLE cars_by_cost ADD cc INT;`
- `ALTER TABLE cars_by_cost  
ALTER cc TYPE BIGINT;`
  - Types must be compatible
- `ALTER TABLE cars_by_cost DROP cc;`



## CQL: Remove or Empty Table


---

- To fully remove a Table:

```
DROP TABLE cars_by_cost;
```

- To clear all data (delete all Partitions) from a Table – but spare the Table:

```
TRUNCATE cars_by_cost;
```



## CQL: Read Data (I)

---

- General syntax:

**SELECT columns**

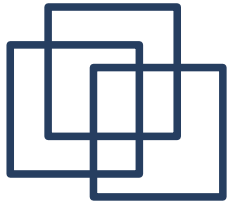
**FROM table**

WHERE relations

ORDER BY clustering\_column ASC/DESC

LIMIT number

ALLOW FILTERING;

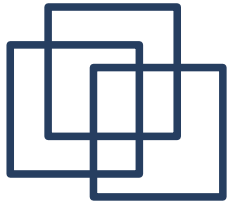


## CQL: Read Data (II)

---

- Typical cases
  - Beware, these examples include Anti-Patterns!
- `SELECT brand, merchant`  
`FROM cars_by_cost;`
  - Avoid retrieving all partitions and rows unless absolutely necessary
- `SELECT *`  
`FROM cars_by_cost;`
  - Avoid retrieving all columns unless necessary

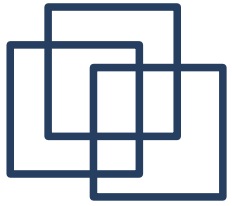




## CQL: Read Data (III)

---


- `SELECT * FROM cars_by_cost  
WHERE brand = “b” AND model = “m”;`
  - To retrieve a partition, values for **all** partition columns are needed
- `SELECT * FROM cars_by_cost  
WHERE brand = “b” AND  
model = “m” AND cost < 1000;`
  - To retrieve a row, values for **all** partition and clustering columns (primary key) are needed



## CQL: Read Data (IV)

---

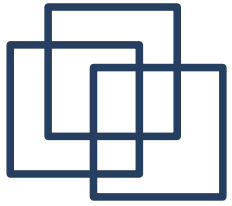
- Secondary Indexes allow to query normal columns
  - Their usage is NOT a spontaneous decision, but a well thought one
- **CREATE INDEX** merchant\_idx  
ON cars\_by\_cost (**merchant**);
  - SELECT \* FROM cars\_by\_cost  
WHERE **merchant** = “m”;
  - SELECT \* FROM cars\_by\_cost  
WHERE brand = “b” AND **merchant** = “m”;
- **DROP INDEX** merchant\_idx;



## CQL: Read Data (V)

---

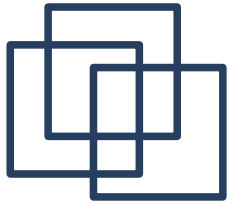
- Trick of ALLOW FILTERING
  - Allows scanning over all partitions and the predicate needs not give values for all partition columns
  - May lead to slow queries with large result set



# CQL: Additional Functions

---

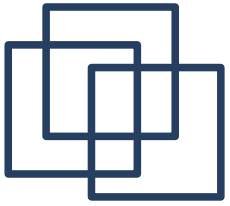
- Aggregation related
  - count(), min(), max(), sum(), avg(), ...
- Time related
  - now(), dateof(), ...
- Blob conversion related
  - bigintAsBlob, blobAsBigint, ...
- User Defined Functions are also possible!
  - To be executed within C\*, thus written in Java



## CQL: Create (Insert) Data

---

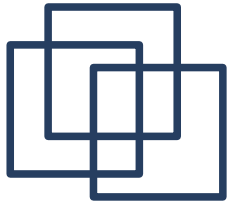
- INSERT INTO  
cars\_by\_cost (brand, model, cost, merchant)  
VALUES (“volvo”, “xc90”, 9999, “daves”);
- What does it do?
  - Creates non-existing partitions
  - **But also updates existing partitions**



# CQL: Update Data

---


- UPDATE cars\_by\_cost  
SET merchant = “pauls”  
WHERE brand = “volvo”  
AND model = “xc90”  
AND cost = 9999;
- What does it do?
  - Updates existing partitions
  - **But also creates non-existing partitions**



# CQL: Upsert Data

---

- Insert and Update have the notion of Upsert
  - Update or Insert
- Why?
  - Because of the way data is organized into Clustering columns



## CQL: Delete Data (I)

---

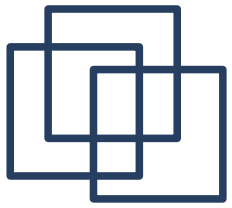
- Deleting a Partition

```
DELETE FROM cars_by_cost  
WHERE brand = “b” AND model = “m”;
```

- Deleting a Row

```
DELETE FROM cars_by_cost  
WHERE brand = “b”  
      AND model = “m”  
      AND cost = 1000;
```





## CQL: Delete Data (II)


---

- Deleting (setting to NULL) a cell from a Row


```
DELETE merchant FROM cars_by_cost  
WHERE brand = "b"  
      AND model = "m"  
      AND cost = 1000;
```

- To clear all data (delete all Partitions) from a Table – but spare the Table:

```
TRUNCATE cars_by_cost;
```



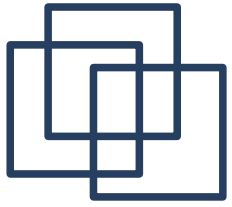
# ACID & TRANSACTIONS



# ACID

---

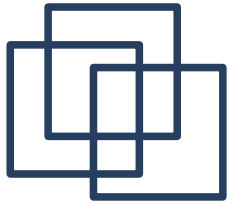
- Not in the usual RDBMS sense
- Atomicity
  - Per Partition
- Consistency
  - Configurable via CL
- Isolation
  - Per Partition
- Durability
  - Write operations are indeed persisted



# Lightweight Transactions

---

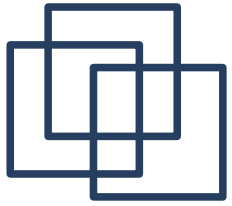
- Two ways to accomplish, as
  - Compare-And-Set (CAS) operations
  - Batch Statements
- Both affect performance
  - Negative impact



## Compare-And-Set Ops. (I)

---

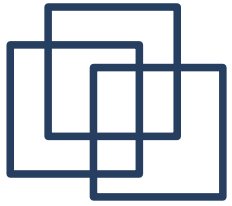
- It performs a Read operation, checks a Condition, and if that one holds, proceeds with the Write operation
  - All atomically



## Compare-And-Set Ops. (II)

---


- INSERT INTO  
cars\_by\_cost (brand, model, cost, merchant)  
VALUES (“volvo”, “xc90”, 9999, “daves”)  
**IF NOT EXISTS;**
- UPDATE cars\_by\_cost  
SET merchant = “pauls”  
WHERE brand = “volvo”  
AND model = “xc90”  
AND cost = 9999  
**IF EXISTS;**



# Batch Statements (I)

---

- BATCH statement
- Offers Atomicity for a series of operations
  - Write-Operations
    - INSERT, UPDATE, DELETE
  - All these operations receive the same timestamp
  - Order of operations is NOT guaranteed
- Does NOT offer isolation
  - Other statements can read/write data affected by the batch




## Batch Statements (II)

---


### BEGIN BATCH

```
UPDATE user SET lock = true IF lock = false;  
WHERE performer = 'The Beatles' AND  
           year = 1966 AND title = 'Revolver';  
INSERT INTO albums_by_performer (performer, year, title,  
                                genre)  
VALUES ('The Beatles', 1966, 'Revolver', 'Rock');  
UPDATE user SET lock = false;
```

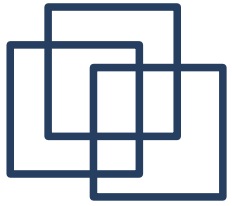
### APPLY BATCH;








# DATA MODELING FOR C\*



# Data Modeling (DM)

---

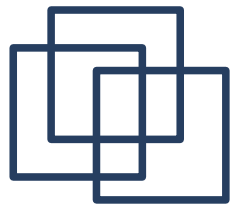
1. Conceptual Data Model
2. Query-Driven Schema Design
  - Access Patterns
3. Logical Data Model
4. Analysis for Partition Size and Data Duplication
5. Physical Data Model
  - CQL



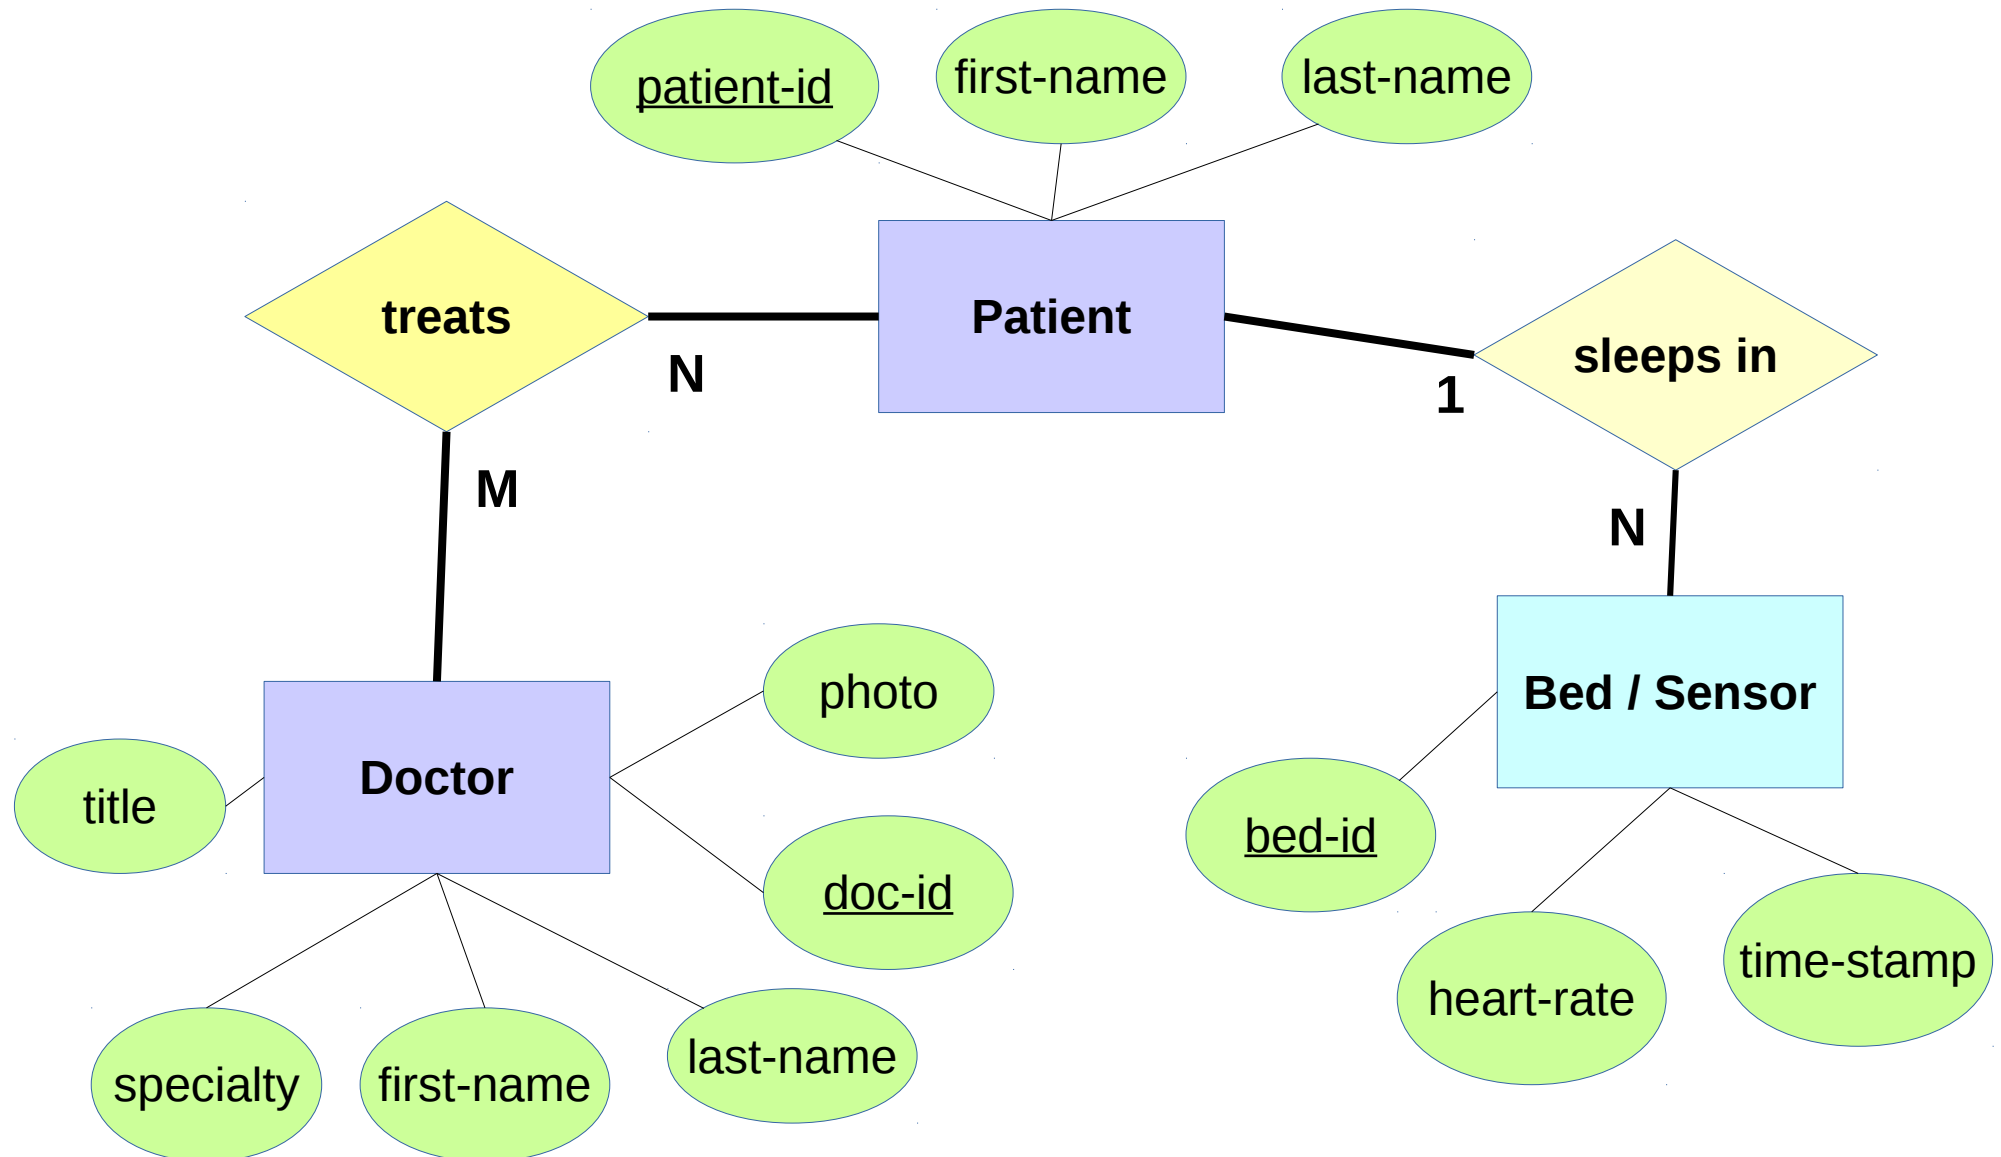
# DM: Example

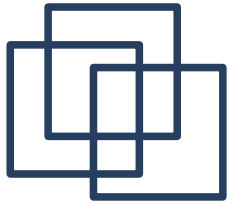
---

- Toy Application related to heart rate measurements of patients in a hospital
- We start with a Conceptual Model and some Queries in the form of Requirements



# DM: Conceptual Data Model






## DM: Queries – Reqs. (I)

---

- Q1) Retrieve all information for a Doctor given his/her full name
- Q2) Retrieve all information excluding the photograph for a Doctor given his/her full name
- Q3) Retrieve the names and ids of the Doctors treating a given Patient, who is known by his/her patient-id. Sort them alphabetically.



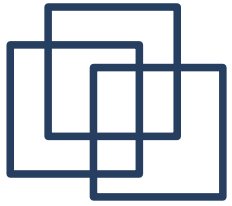
## DM: Queries – Reqs. (II)

---

Q4) Find the average heart rate for a given Patient – known by his/her patient-id – on a single given date and given hour range

Q5) Find the average heart rate for a given Patient – known by his/her patient-id – on a given date range.

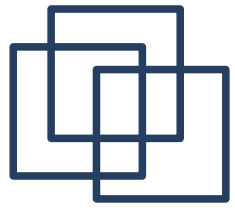
Assume that the date range will have at most ten days.



# DM: Logical Data Model (I)

---

- Q1, Q2:
  - last\_name TEXT → Partition Key column (\*)
  - first\_name TEXT → Partition Key column (\*)
  - doc\_id INT → Clustering column (\*\*)
  - specialty TEXT
  - title TEXT
  - photo BLOB
- (\*) We search by it
- (\*\*) Needed for uniqueness across a row - case of two different doctors with same names

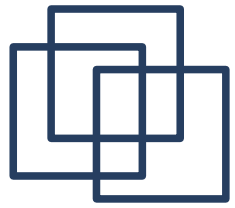


## DM: Logical Data Model (II)

---

- Q3:
  - patient\_id        TEXT    → Partition Key column (\*)
  - doc\_last\_name    TEXT    → Clustering column (\*\*)
  - doc\_first\_name   TEXT    → Clustering column (\*\*)
  - doc\_id            INT     → Clustering column (\*\*\*)
- (\*) We search by it
- (\*\*) We order the result by them
- (\*\*\*) Needed for uniqueness across a row - Corner case of two different doctors with same names treating one patient






## DM: Logical Data Model (III)

---

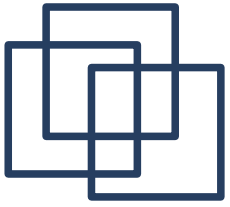
- Q4, Q5:
  - patient\_id TEXT → Partition Key column (\*)
  - patient\_last\_name TEXT
  - patient\_first\_name TEXT
  - bed\_id TEXT
  - when TIMESTAMP → Clustering column (\*\*)
  - heart\_rate INT
- (\*) We search by it
- (\*\*) We perform range-search by it



# DM: Analysis (I)

---

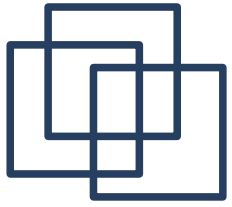
- Q1, Q2:
  - A photo is only needed for Q1 but can slow down Q2
  - We can duplicate data to make Q2 faster
- Q1 =  
[last\_name **K**, first\_name **K**, doc-id **C** ↑,  
specialty, title, **photo** ]
- Q2 =  
[last\_name **K**, first\_name **K**, doc\_id **C** ↑,  
specialty, title ]



## DM: Analysis (II)

---


- Q3: No surprises
- Q3 =  
[patient\_id **K**,  
doc\_last\_name **C** ↑,  
doc\_first\_name **C** ↑,  
doc\_id **C** ↑]



## DM: Analysis (III)

---

- Q4, Q5:
  - We have at most 1 measurement per minute, that is 1440 per day, 14 400 for ten days
  - The range is defined by the timestamp is quite large
  - Although we keep a full timestamp, we only query for days or hours
  - Information about the patient's names is repeated



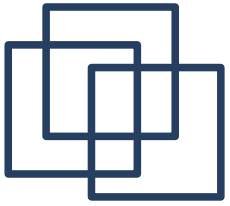
## DM: Analysis (IV)

---

- Q4, Q5 =

```
[patient_id      K,  
patient_last_name S,  
patient_first_name S,  
bed_id,  
when_date      TIMESTAMP C↓,  
when_day_minutes INT      C↓,  
heart_rate ]
```


- This is good enough for bounded Partition size



## DM: Physical and CQL

---

- Q1) TABLE docs\_w\_photos\_by\_name
- Q2) TABLE docs\_by\_name
- Q3) TABLE docs\_by\_patient
- Q4, Q5)  
TABLE heart\_rate\_by\_patient\_and\_time




# DM: CQL and Application

---




**Show me some Code**



---

# MORE ABOUT C\*

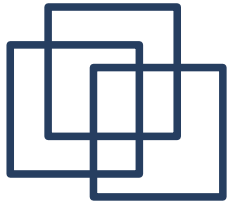




# Get Started (I)

---

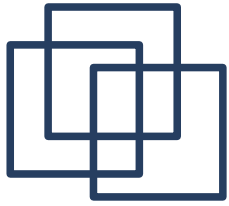
- Apache C\* - <http://cassandra.apache.org/>
- DataStax C\* - <http://www.datastax.com/>
- DataStax C\* Drivers and Tools - <https://academy.datastax.com/downloads/welcome>
- DataStax and Apache C\* Drivers -
  - DataStax
  - Apache



## Get Started (II)

---

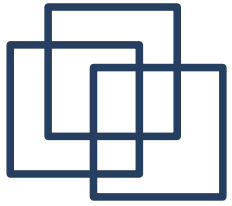
- DataStax and Apache CQL Documentation -
  - <http://docs.datastax.com/en/cql/3.3/cql/cqlIntro.html>
  - <https://cassandra.apache.org/doc/cql/CQL.html>
- **DataStax Startup Program** -  
<http://www.datastax.com/datastax-enterprise-for-startups>



# Learn More

---

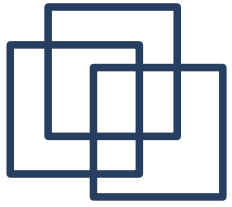
- DataStax Community Service
  - <http://www.planetcassandra.org/>
- DataStax Cassandra Academy
  - <https://academy.datastax.com/>
- Stack Overflow Tags:
  - [cassandra](#), [datastax](#), [datastax-enterprise](#)
- Books
  - visit [Amazon](#)



# Get Involved

---

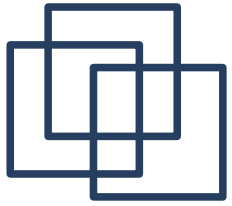
- Events about Cassandra - <http://www.datastax.com/company/events>
  - Conferences
  - Webinars
- Meetup Groups - <http://www.meetup.com/>
  - Join a group, or
  - **CREATE ONE!**



## References – Sources (I)

---

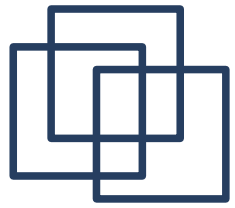
- These (already mentioned) web pages:
  - <http://cassandra.apache.org/>
  - <http://www.datastax.com/>
  - <http://docs.datastax.com/en/cql/3.3/cql/cqlIntro.html>
  - <https://cassandra.apache.org/doc/cql/CQL.html>
  - <https://academy.datastax.com/>
  - <http://www.planetcassandra.org/>



## References – Sources (II)

---


- These additional articles/web pages:
  - Leslie Lamport, "Time, clocks, and the ordering of events in a distributed system",  
[http://research.microsoft.com/en-us/um/people/la  
mport/pubs/time-clocks.pdf](http://research.microsoft.com/en-us/um/people/la<br/>mport/pubs/time-clocks.pdf)
  - Mark Burgess, "Deconstructing the 'CAP theorem' for CM and DevOps",  
[http://markburgess.org/blog\\_cap.html](http://markburgess.org/blog_cap.html)



# References - Acknowledgments

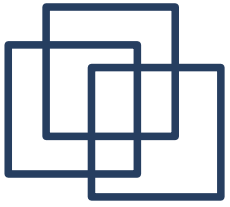
---

- Tables and Graphics in slides marked with





are copyrighted work of **DATASTAX**  
and used here after their permission.

- We really thank them for that.




# Let's Talk!

---







---

# THANK YOU

